Phase 5: Semantic Expert Diagnosis

Deep Contextual Representation & Gradient-Based Token Highlights

Model: Fine-Tuned DistilRoBERTa / MuRIL Transformer
P(FAKE): 95.0%
Token Attributions (Integrated Gradients):